

Visualization in bioinformatics

Abhijit Dasgupta, PhD

Networks

Visualizing a proteomic network

We read a dataset that contains the network relationships between different proteins

```
library(ggnetwork)
datf <- rio::import('data/string_graph.txt')
head(datf)
```

	node1	node2	node1_string_id	node2_string_id	node1_external_id
1	CXCR3	CCR7	1855969	1843829	ENSP00000362795
2	ITGA4	EED	1858446	1845338	ENSP00000380227
3	SMC3	CENPK	1854200	1843648	ENSP00000354720
4	HNRNPA1	LUC7L3	1852510	1843556	ENSP00000341826
5	SMC2	RB1	1847012	1845924	ENSP00000286398
6	RBBP4	CENPK	1855919	1843648	ENSP00000362592

	node2_external_id	neighborhood	fusion	cooccurence	homology	coexpression
1	ENSP00000246657	0	0	0	0.847	0.000
2	ENSP00000263360	0	0	0	0.000	0.000
3	ENSP00000242872	0	0	0	0.000	0.000
4	ENSP00000240304	0	0	0	0.000	0.000
5	ENSP00000267163	0	0	0	0.000	0.136
6	ENSP00000242872	0	0	0	0.000	0.000

	experimental	knowledge	textmining	combined_score
1	0.000	0.9	0.878	0.913
2	0.566	0.0	0.312	0.688
3	0.000	0.9	0.081	0.904

Visualizing a proteomic network

The **igraph** package allows the creation of network graphs.

However, here, we're only using it for data ingestion

```
pacman::p_load(igraph)
grs <- graph_from_data_frame(datf[,c('node1', 'node2')],
                             directed = F)
grs
```


```
IGRAPH 473d753 UN-- 37 58 --
+ attr: name (v/c)
+ edges from 473d753 (vertex names):
[1] CXCR3 --CCR7 ITGA4 --EED SMC3 --CENPK HNRNPA1--LUC7L3
[5] SMC2 --RB1 RBBP4 --CENPK CXCR5 --CXCL13 CD44 --RUNX2
[9] CXCR5 --PF4 PF4 --THBD SMARCA1--EZH2 HMMR --BARD1
[13] MBP --MMP7 CCL19 --CCR7 RBBP4 --EZH2 RUNX2 --RB1
[17] RB1 --HSPA8 DHX9 --BARD1 CXCL13 --CCR7 SMC2 --KIF23
[21] CD44 --HMMR ITGA4 --CD44 RB1 --SMARCE1 ITGA4 --CCR7
[25] MBP --STK4 RBBP4 --LIN9 RB1 --EED CXCR5 --CCR7
[29] PSMA1 --HSPA8 RBBP4 --SMARCA1 CXCR3 --ITGA4 MBP --CDK12
+ ... omitted several edges
```

We see that this object holds the different connections.

Visualizing a proteomic network

We can then transform this data into ggplot-friendly data, to use ggplot for the plotting

```
library(intergraph)
ggdf <- ggnetwork(asNetwork(gr),
                  layout='fruchtermanreingold')
ggplot(ggdf, aes(x = x, y = y,
                 xend = xend, yend = yend)) +
  geom_edges(color = "black",
             curvature = 0.1,
             size = 0.95, alpha = 0.8) +
  geom_nodes(aes(x = x, y = y),
             size = 3,
             alpha = 0.5,
             color = "orange") +
  geom_nodelabel_repel(aes(label = vertex.names),
                      size=4, color="#8856a7") +
  theme_blank() + theme(legend.position = "none")
```



Composing different genomic data into tracks

The ggbio package

The **ggbio** package has several functions that allow graphical representations of different genomic entities.

You will see a lot of use of `autoplot`, which is a software technique to create default visualizations based on the type of entry.

An ideogram


```
library(ggbio) # p_install('ggbio', try.bioconductor=  
p.ideo <- Ideogram(genome = 'hg19')  
p.ideo
```



The ggbio package

Visualizing the gene model

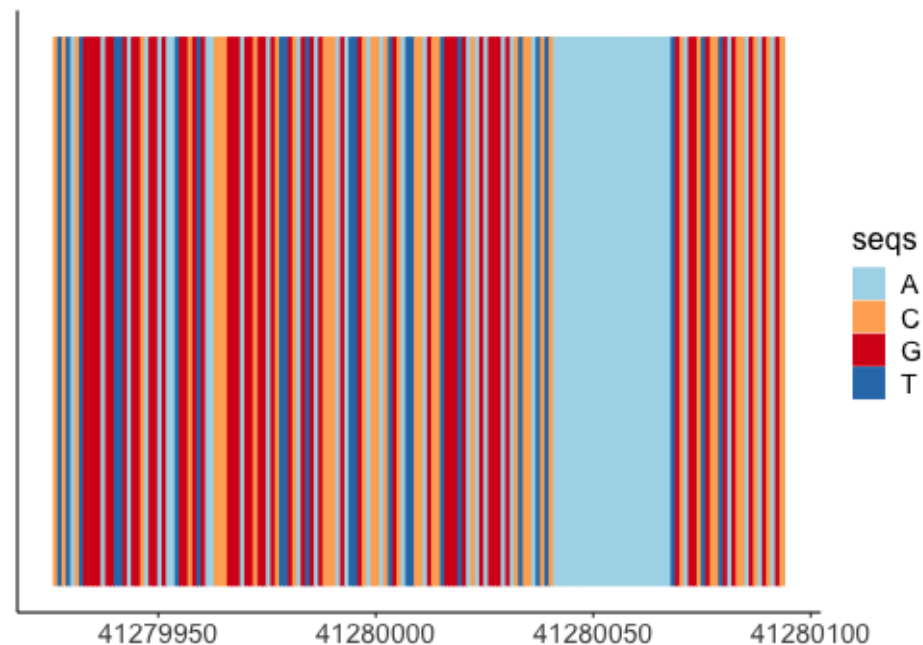
```
pacman::p_load(Homo.sapiens)
data(genesymbol, package='biovizBase')
wh <- genesymbol[c('BRCA1','NBR1')]
wh <- range(wh, ignore.strand=T)
p.txdb <- autoplot(Homo.sapiens, which = wh)
p.txdb
```



The ggbio package

A reference track

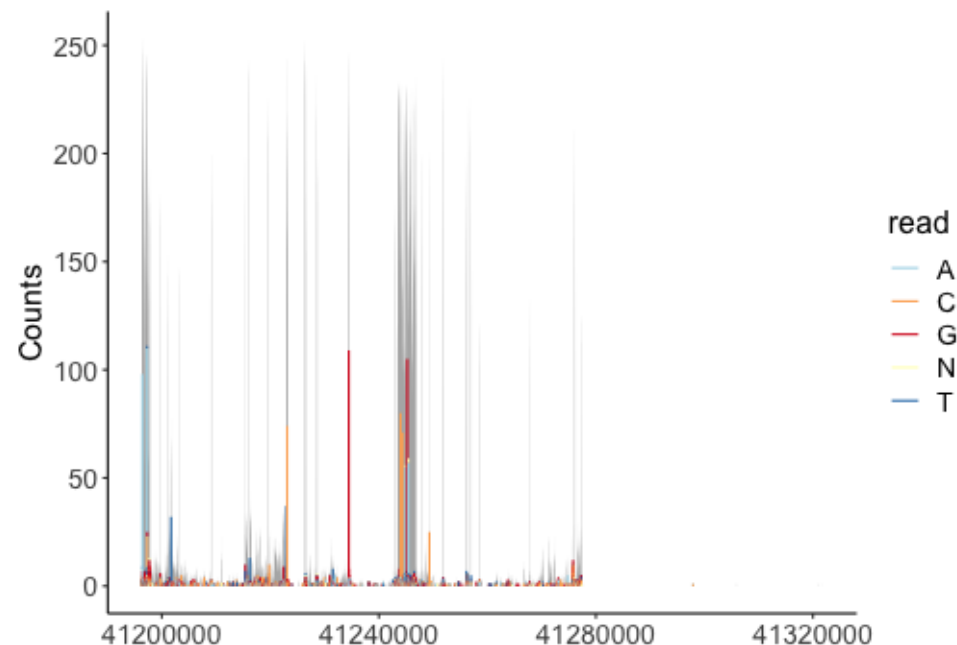
```
library(BSgenome.Hsapiens.UCSC.hg19)
bg <- BSgenome.Hsapiens.UCSC.hg19
p.bg <- autoplot(bg, which=wh)
p.bg + zoom(1/1000)
```



The ggbio package

An alignment track with mismatch proportions


```
library(BSgenome.Hsapiens.UCSC.hg19)
fl.bam <- system.file("extdata", "wg-brca1.sorted.bam", package="BSgenome.Hsapiens.UCSC.hg19")
wh <- keepSeqlevels(wh, "chr17")
bg <- BSgenome.Hsapiens.UCSC.hg19
p.mis <- autoplot(fl.bam, bsgenome = bg, which = wh,
p.mis
```



The ggbio package

Putting it into tracks


```
pacman::p_load(GenomicRanges)
gr17 <- GRanges("chr17", IRanges(41234415, 41234569))
tk <- tracks(p.ideo,
             ref=p.bg,
             mismatch=p.mis,
             gene=p.txdb,
             heights=c(2,1,3,4))+
  xlim(gr17) +
  theme_tracks_sunset()
print(tk)
```



The ggbio package

A karyogram

```
data(ideoCyto, package = "biovizBase")
autoplot(ideoCyto$hg19, layout = "karyogram",
         cytobands = TRUE)
```



P-values and Manhattan plots

A very simple example

```
library(tidyverse)
clinical <- rio::import('data/BreastCancer_Clinical.xlsx') %>% janitor::clean_names()
proteome <- rio::import('data/BreastCancer_Expression.xlsx') %>% janitor::clean_names()
final_data <- clinical %>%
  inner_join(proteome, by = c('complete_tcga_id' = 'tcga_id')) %>%
  dplyr::filter(gender == 'FEMALE') %>%
  dplyr::select(complete_tcga_id, age_at_initial_pathologic_diagnosis, er_status, starts_with("np"))
head(final_data)
```

	complete_tcga_id	age_at_initial_pathologic_diagnosis	er_status	np_958782	np_958785	np_958786	np_000436
1	TCGA-A2-A0CM		40 Negative	0.6834035	0.6944241	0.6980976	0.6870771
2	TCGA-BH-A18Q		56 Negative	0.1953407	0.2154129	0.2154129	0.2053768
3	TCGA-A7-A0CE		57 Negative	-1.1231731	-1.1231731	-1.1168605	-1.1294857
4	TCGA-D8-A142		74 Negative	0.5385958	0.5422105	0.5422105	0.5349810
5	TCGA-A0-A0J6		61 Negative	0.8311317	0.8565398	0.8565398	0.8367780
6	TCGA-A2-A0YM		67 Negative	0.6558497	0.6581426	0.6558497	0.6558497
	np_958781	np_958780	np_958783	np_958784	np_112598	np_001611	
1	0.6870771	0.6980976	0.6980976	0.6980976	-2.6521501	-0.9843733	
2	0.2154129	0.2154129	0.2154129	0.2154129	-1.0357599	-0.5172257	
3	-1.1294857	-1.1200168	-1.1231731	-1.1231731	2.2445844	-2.5750648	
4	0.5422105	0.5422105	0.5422105	0.5422105	-0.1482049	0.2674902	
5	0.8650092	0.8565398	0.8508936	0.8508936	-0.9671961	2.8383705	
6	0.6512639	0.6581426	0.6558497	0.6558497	-1.9695337	1.3070365	

A very simple example

```
results <- final_data %>%
  summarise_at(vars(starts_with('np')),
    ~wilcox.test(. ~ er_status)$p.value)
results
```

```
      np_958782 np_958785 np_958786 np_000436 np_958781 np_958780 np_958783 np_958784 np_112598      np_001611
1 0.6988415 0.6910103 0.6832121 0.6910103 0.6832121 0.6910103 0.6910103 0.6832121 0.9957714 0.0001218627
```

. is the placeholder for what's specified inside the vars().

This isn't in the right format for me to plot

A very simple example


```
results %>% tidyr::pivot_longer(cols=everything(),  
                                names_to='protein',  
                                values_to='pvalue')
```

```
# A tibble: 10 x 2  
  protein      pvalue  
  <chr>      <dbl>  
1 np_958782 0.699  
2 np_958785 0.691  
3 np_958786 0.683  
4 np_000436 0.691  
5 np_958781 0.683  
6 np_958780 0.691  
7 np_958783 0.691  
8 np_958784 0.683  
9 np_112598 0.996  
10 np_001611 0.000122
```


A very simple example


```
theme_439 <- theme_bw() +  
  theme(axis.title = element_text(size=16),  
        axis.text = element_text(size=8))
```

```
results %>% pivot_longer(  
  cols=everything(),  
  names_to='protein',  
  values_to='pvalue') %>%  
  ggplot(aes(x = protein, y = pvalue)) +  
  geom_point() +  
  theme_439
```




A very simple example

```
pacman::p_load('ggalt')
results %>% pivot_longer(
  cols=everything(),
  names_to = 'protein',
  values_to = 'pvalue') %>%
  ggplot(aes(x = protein, y = pvalue)) +
  geom_point() +
  geom_lollipop(point.colour='blue', point.size=4)+
  coord_flip()+
  theme_439
```



A very simple example

```
results %>% pivot_longer(
  cols=everything(),
  names_to = 'protein',
  values_to = 'pvalue') %>%
  ggplot(aes(x = protein, y = pvalue)) +
    geom_point() +
    geom_lollipop(point.colour='blue', point.size=4)
    geom_hline(yintercept = 0.05, linetype=2)+
    coord_flip() +
    theme_439
```




Manhattan plot

A Manhattan plot is used to visualize a set of p-values from unit-based tests

It plots the negative log p-value at each unit

```
results %>% pivot_longer(
  cols=everything(),
  names_to = 'protein',
  values_to = 'pvalue') %>%
  ggplot(aes(x = protein, y = -log10(pvalue))) +
    geom_point() +
    geom_lollipop() +
    geom_hline(yintercept = 8, linetype=2)+
    labs(x = 'Protein',
         y = expression(log[10](p-value))) +
    theme_439
```



Manhattan plot


There is a specialized package for doing Manhattan plots and quantile plots for GWAS data

This package is meant to work with PLINK output, but the function is generic

```
library(qqman)  
manhattan(gwasResults)
```

Manhattan plot

```
library(qqman)
manhattan(gwasResults,
          annotatePval = 1e-6,
          annotateTop=F)
```



Heatmaps

Let us count the ways

There are several ways of doing heatmaps in R:

- <https://jokergoo.github.io/ComplexHeatmap-reference/book/>
- http://sebastianraschka.com/Articles/heatmaps_in_r.html
- <https://plot.ly/r/heatmaps/>
- <http://moderndata.plot.ly/interactive-heat-maps-for-r/>
- <http://www.siliconcreek.net/r/simple-heatmap-in-r-with-ggplot2>
- <https://rud.is/b/2016/02/14/making-faceted-heatmaps-with-ggplot2/>

Some example data


```
library(Biobase)
#data(sample.ExpressionSet)
exdat <- readRDS('data/exprset.rds')
library(limma)
design1 <- model.matrix(~type, data=pData(exdat))
lm1 <- lmFit(exprs(exdat), design1)
lm1 <- eBayes(lm1) # compute linear model for each probeset
geneID <- rownames(topTable(lm1, coef = 2, number = 100,
                           adjust.method = 'none',
                           p.value = 0.05))
exdat2 <- exdat[geneID,] # Keep features with p-values < 0.05

head(exdat2)
```

```
ExpressionSet (storageMode: lockedEnvironment)
assayData: 1 features, 26 samples
  element names: exprs, se.exprs
protocolData: none
phenoData
  sampleNames: A B ... Z (26 total)
  varLabels: sex type score
  varMetadata: labelDescription
featureData: none
experimentData: use 'experimentData(object)'
Annotation: hgu95av2
```

Using Heatplus

```
# BiocManager::install('Heatplus')
library(Heatplus)
reg1 <- regHeatmap(exprs(exdat2), legend=2, col=heat.
                    breaks=-3:3)
plot(reg1)
```




Using Heatplus

```

corrdist <- function(x) as.dist(1-cor(t(x)))
hclust.av1 <- function(x) hclust(x, method='average')
reg2 <- regHeatmap(exprs(exdat2), legend=2,
                   col=heat.colors,
                   breaks=-3:3,
                   dendrogram =
                     list(clustfun=hclust.av1,
                          distfun=corrdist))

plot(reg2)

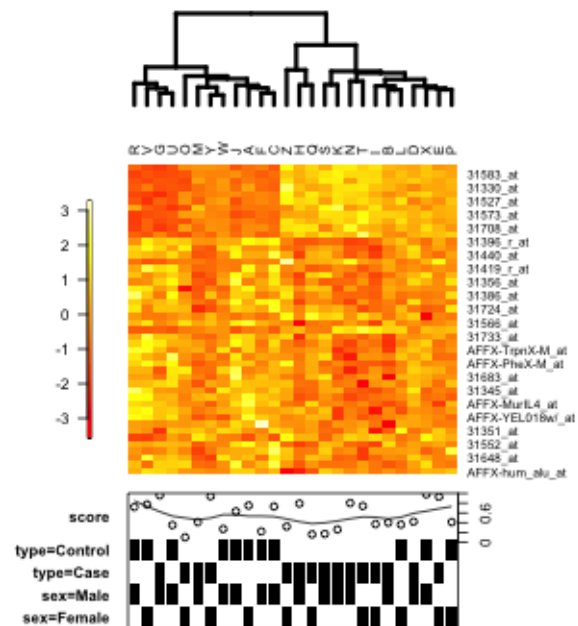
```



Using Heatplus

Adding annotations

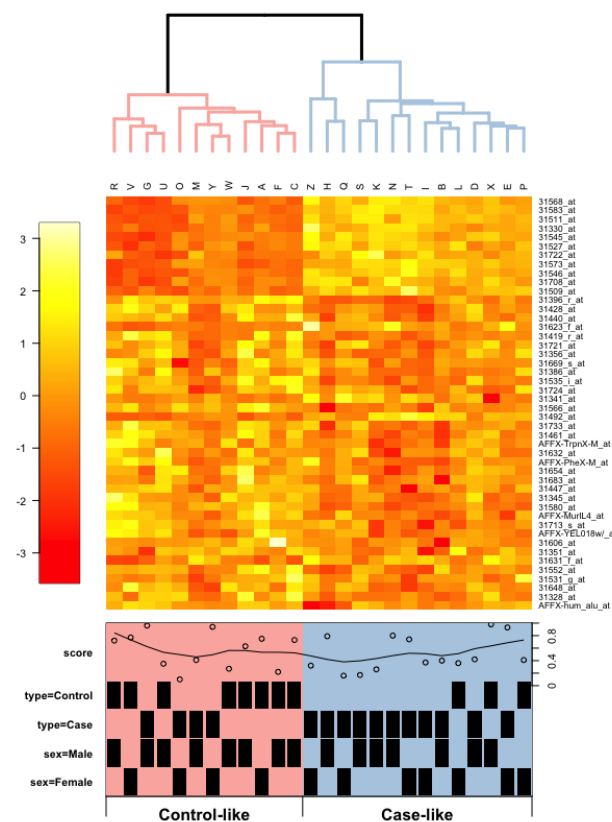
```
ann1 <- annHeatmap(exprs(exdat2),
  ann=pData(exdat2),
  col = heat.colors)
plot(ann1)
```



Using Heatplus


Adding annotations

```
ann2 <- annHeatmap(exprs(exdat2),
                    ann=pData(exdat2),
                    col = heat.colors,
                    cluster =
                      list(cuth=7500,
                           label=c('Control-like', 'Case-like')))
plot(ann2)
```



Using ComplexHeatmap


Source code [here](#)




UpSet plots

UpSet plots are nice visualizations for looking at commonalities (complex intersections) between sets of objects.

We used UpSet plots to look at missing value patterns

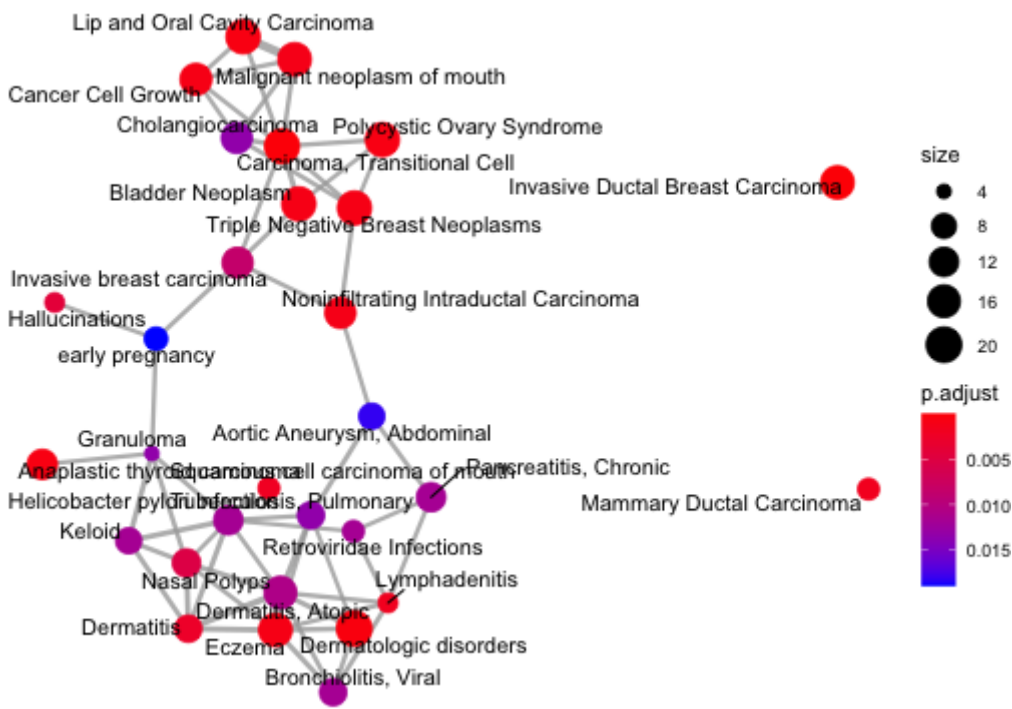


UpSet plots



clusterProfiler

Enrichment network based on GSEA



Playing with Seurat

Example data

```
library(Seurat)
# pbmc.data <- Read10X(data.dir='data/hg19/')
# pbmc <- CreateSeuratObject(counts = pbmc.data, project='pbmc3k', min.cells=3, min.features=200)
pbmc <- readRDS('data/pbmc.rds')
pbmc
```

An object of class Seurat
 13714 features across 2700 samples within 1 assay
 Active assay: RNA (13714 features, 0 variable features)

```
names(pbmc)
```

```
[1] "RNA"
```

```
slotNames(pbmc)
```

```
[1] "assays"      "meta.data"   "active.assay" "active.ident" "graphs"      "neighbors"   "reductions"
[8] "project.name" "misc"        "version"      "commands"    "tools"
```

Adding QC metrics and plotting

We'll calculate mitochondrial QC metrics (percentage counts originating from mitochondrial genes)

```
pbmc[['percent.mt']] <- PercentageFeatureSet(pbmc, pattern = '^MT-')  
head(pbmc@meta.data)
```


	orig.ident	nCount_RNA	nFeature_RNA	percent.mt
AAACATACAACCAC	pbmc3k	2419	779	3.0177759
AAACATTGAGCTAC	pbmc3k	4903	1352	3.7935958
AAACATTGATCAGC	pbmc3k	3147	1129	0.8897363
AAACCGTGCTTCCG	pbmc3k	2639	960	1.7430845
AAACCGTGATGCG	pbmc3k	980	521	1.2244898
AAACGCACTGGTAC	pbmc3k	2163	781	1.6643551

Visualizing metrics

```
# plt <- VlnPlot(object = pbmc,
#   features = c('nFeature_RNA',
#                 'nCount_RNA',
#                 'percent.mt'))


plot_data <- pbmc@meta.data %>%
  tidyr::gather(variable, value, -orig.ident)

ggplot(plot_data, aes(orig.ident, value)) +
  geom_violin(fill = 'red') +
  geom_jitter(width=0.5, alpha = 0.1) +
  facet_wrap(~variable, nrow = 1,
             scales = 'free_y') +
  labs(x = 'Identity', y = 'Expression Level') +
  theme_classic()
```



Visualizing feature-feature relationships

```
plot1 <- FeatureScatter(object = pbmc,  
                        feature1 = "nCount_RNA",  
                        feature2 = "percent.mt")  
plot2 <- FeatureScatter(object = pbmc,  
                        feature1 = "nCount_RNA",  
                        feature2 = "nFeature_RNA")  
CombinePlots(plots = list(plot1, plot2))
```



Visualizing feature-feature relationships


```

cormatrix <- cor(pbmc@meta.data %>% dplyr::select(-or
plt1 <-
  ggplot(pbmc@meta.data,
    aes(x = nCount_RNA,
        y = percent.mt,
        group = orig.ident,
        color = orig.ident)) +
  geom_point() +
  theme_classic() +
  labs(color = 'Identity',
        title=as.character(round(cormatrix['nCount_R
  theme(plot.title = element_text(face = 'bold', hjus

plt2 <-
  ggplot(pbmc@meta.data,
    aes(x = nCount_RNA,
        y = nFeature_RNA,
        group = orig.ident,
        color = orig.ident)) +
  geom_point() +
  theme_classic() +
  labs(color = 'Identity',
        title=as.character(round(cormatrix['nCount_RNA
  theme(plot.title = element_text(face = 'bold', hjus

ggpubr::ggarrange(plt1, plt2, nrow = 1, ncol=2)

```




Feature selection

```
pbmc <- subset(x = pbmc,
  subset = nFeature_RNA > 200 & nFeature_RNA < 2500
pbmc <- NormalizeData(object = pbmc,
  normalization.method = "LogNorm
  scale.factor = 10000)
# This is stored in pbmc[['RNA']]@meta.features

pbmc <- FindVariableFeatures(object = pbmc,
  selection.method = "vst"
  nfeatures = 2000)

# Identify the 10 most highly variable genes
top10 <- head(x = VariableFeatures(object = pbmc), 10)

# plot variable features with and without labels
plot1 <- VariableFeaturePlot(object = pbmc)
plot1
```

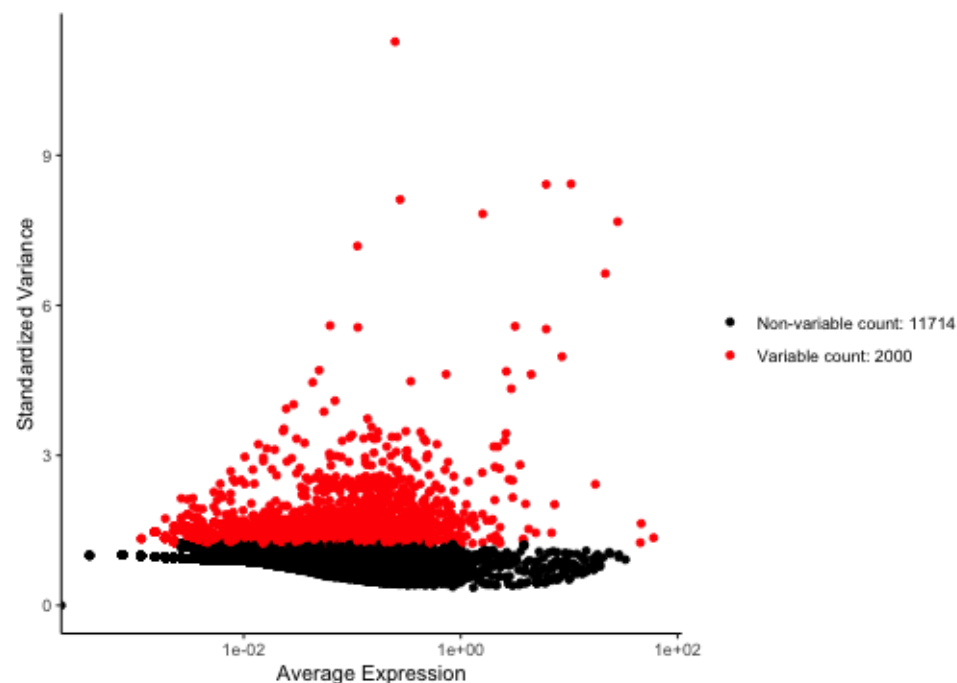


Feature selection

```
plt_data <- pbmc[['RNA']]@meta.features %>%
  rownames_to_column(var='id')
topvars <- pbmc[['RNA']]@var.features
plt_data <- plt_data %>%
  mutate(indic = ifelse(id %in% topvars,
                        'Variable count',
                        'Non-variable count'))


bl <- plt_data %>%
  dplyr::count(indic) %>%
  glue::glue_data("{indic}: {n}")
names(bl) <- c('Non-variable count', 'Variable count')
plt_data <- plt_data %>%
  mutate(indic = bl[indic])
plt11 <- ggplot(plt_data,
               aes(x = vst.mean,
                   y = vst.variance.standardized,
                   color = indic)) +

  geom_point() +
  scale_x_log10() +
  scale_color_manual(values = c('black', 'red')) +
  labs(x = 'Average Expression', y = 'Standardized Va
  theme_classic()
plt11
```



Feature selection

```
# plot2 <- LabelPoints(plot = plot1, points = top10,  
plt12 <- plt11 + ggrepel::geom_text_repel(data = plt_  
aes(label =  
color = 'bl  
  
plt12
```



There's a lot more

We'll stop our sampling here.

- Many Bioconductor packages do use ggplot, however some use base graphics
 - Faster
- Key is to find where the data is stored, and use that to create visualizations
- Bioconductor tends to create
 - One monolithic object
 - Containing different information in slots
 - combined by lists
- `slotNames` and `names` are your friends